

Training Lifecycle

How Raw Text Becomes an Intelligent Model

The Big Question: You now understand tokens, embeddings, and the Transformer architecture. But how does a model actually **learn**? What happens between raw internet text and a model that can write code, answer questions, and reason through problems? The answer is the **three-stage training lifecycle**: Pre-Training → SFT → Alignment.

The Three-Stage Training Pipeline

Stage	Name	Objective	Duration	Cost
1	Pre-Training	Next-token prediction on trillions of tokens; learns language, facts, reasoning	Months	\$10M–\$100M+
2	Supervised Fine-Tuning (SFT)	Instruction-response pairs; transforms completion model into helpful assistant	Days–weeks	\$10K–\$1M
3	Alignment (RLHF / DPO)	Human preferences shape outputs; safer, more helpful, less harmful	Days	\$1K–\$100K

Most companies skip Stage 1 entirely. Google, Microsoft, Amazon, Netflix — they all start from a pre-trained base model (LLaMA-3, Mistral, Claude, GPT-4) and apply Stages 2 and 3 for their specific use case. Pre-training is only done by a handful of labs with \$100M+ budgets.

Topics Covered Today

#	Topic	Coverage
1	Three-Stage Pipeline	Pre-Training → SFT → Alignment; who does each stage
2	Pre-Training	Next-token prediction, data breakdown, infrastructure at FAANG scale
3	Training Data	CommonCrawl 67%, Books 13%, Code 8%; quality beats quantity
4	Supervised Fine-Tuning (SFT)	Instruction-response pairs; base model vs assistant model

5	SFT Datasets	OpenAssistant, ShareGPT, Alpaca, FLAN, LLaMA-3 internal
6	RLHF	Human preferences → Reward Model → PPO; 3-step alignment
7	DPO	Direct Preference Optimisation; no reward model; used by LLaMA-3, Mistral
8	Constitutional AI	Anthropic's RLAIIF approach; self-critique against principles
9	C = 6ND	Training compute equation; GPT-4 ~\$100M; SFT ~\$32
10	Chinchilla Scaling Laws	Optimal D = 20N; why modern models intentionally over-train
11	FAANG Training	Meta (LLaMA recipe), Google (native multimodal), OpenAI (PRMs), Microsoft (Phi)
12	Architect's Lens	Prompt first, RAG second, fine-tune third; base model selection; cost estimation
13	Hands-On Exercise	SFT dataset inspection + cost calculator with HuggingFace datasets

Section 1 — The Big Picture: Three Stages, One Model

Before a model like GPT-4, LLaMA-3, or Claude reaches you, it goes through three distinct stages. Each has a different objective, dataset, and cost:

```
Raw Internet Text
↓
PRE-TRAINING ← Months, $10M-$100M+, trillions of tokens
(Base Model) ← Knows language, facts, reasoning – but raw/unaligned
↓
SUPERVISED FINE-TUNING (SFT) ← Days/weeks, $10K-$1M
(Instruction Model) ← Knows how to be a helpful assistant
↓
ALIGNMENT (RLHF / DPO) ← Days, $1K-$100K
(Aligned Model) ← Safe, helpful, honest, on-brand
↓
Production LLM (ChatGPT, LLaMA-3-Instruct, Claude, Gemini)
```

Most companies skip Stage 1 entirely. Google, Microsoft, Amazon, Netflix — they all start from a pre-trained base model and apply Stages 2 and 3 for their specific use case. Pre-training is only done by a handful of labs with \$100M+ compute budgets. As an AI engineer, you will spend most of your career in Stages 2 and 3.

Section 2 — Stage 1: Pre-Training

What is pre-training? The model is given a sequence of tokens and asked to predict the next one. That is the entire task.

Input: "The Eiffel Tower is located in"
Target: "Paris"

Input: "def fibonacci(n):\n if n <= 1:\n return"
Target: "n"

Input: "Google was founded in"
Target: "1998"

This single objective — next-token prediction — is repeated across **15 trillion tokens** of text. By the time it is done, the model has learned grammar, syntax, world knowledge, reasoning patterns, code, and multiple languages — purely from predicting the next word.

Pre-training data breakdown (LLaMA-3 style):

Source	Share	Why Included
Web (CommonCrawl)	67%	Sheer volume — diversity of topics, styles, and languages
Books / Literature	13%	Long-form reasoning, narrative structure, coherence over thousands of tokens
GitHub / Code	8%	Structured reasoning, logic, precise syntax — improves reasoning across all tasks
Wikipedia	5%	High-quality factual knowledge, encyclopaedic breadth, clean writing
ArXiv / Papers	4%	Scientific reasoning, formal arguments, domain-specific technical knowledge
Other curated	3%	High-quality niche sources: legal, medical, financial, multilingual

Data quality matters enormously. Microsoft's Phi-3 team showed that training on 3.4T tokens of textbook-quality synthetic data produced a 3.8B model that outperforms 7B models trained on 15T tokens of raw web data. Lesson: garbage in, garbage out — at trillion-token scale.

How pre-training works technically:

1. Tokenize the entire dataset into one giant sequence of token IDs
2. Split into chunks of context_length tokens (e.g. 4,096 tokens each)
3. For each chunk:
 - a. Feed tokens [1..n] into the Transformer
 - b. Predict token [n+1] at every position simultaneously (teacher forcing)
 - c. Compute cross-entropy loss between prediction and actual next token
 - d. Backpropagate gradients through all layers
 - e. Update all parameters with AdamW optimiser
4. Repeat for 15 trillion tokens across thousands of GPUs

Training infrastructure at FAANG scale:

LLaMA-3 70B training (Meta):

Hardware: 16,384 x H100 80GB GPUs

Duration: ~90 days continuous

Framework: PyTorch + FSDP (Fully Sharded Data Parallel)

Throughput: ~1 trillion tokens per day

Total: ~15 trillion tokens

GPT-4 training (estimated):

Hardware: ~25,000 x A100 80GB GPUs

Duration: ~100 days

Cost: \$63M-\$100M (cloud pricing)

Section 3 — Stage 2: Supervised Fine-Tuning (SFT)

The problem with base models: A pre-trained base model is a powerful text predictor — but it will complete your prompt like a web page, not answer it like an assistant.

Prompt: "What is 2+2?"

Base model output: "What is 3+3? What is 4+4? What is 5+5?..."

↑ continuing the pattern of a quiz worksheet – not helpful

SFT model output: "2+2 equals 4."

↑ answers the question directly – helpful

SFT training data format (ChatML — used by OpenAI and LLaMA):

```
<|im_start|>system
You are a helpful assistant.<|im_end|>
<|im_start|>user
Explain gradient descent in simple terms.<|im_end|>
<|im_start|>assistant
Gradient descent is like finding the lowest point in a hilly landscape...<|im_end|>
```

Famous SFT datasets:

Dataset	Size	Source	Used By
OpenAssistant	161K	Human annotators	Open community models
ShareGPT	90K	Real ChatGPT conversations	Vicuna, many others
Alpaca	52K	GPT-4 generated	Stanford Alpaca
FLAN Collection	1.8M	Task-formatted NLP datasets	Google FLAN-T5
LLaMA-3 SFT data	~10M	Meta internal (synthetic + human)	LLaMA-3-Instruct

Key insight: SFT data quality beats quantity every time. 1,000 perfectly crafted examples outperforms 100,000 noisy ones. This is why companies like Google and Meta spend millions on human annotation rather than purely relying on LLM-generated data.

Section 4 — Stage 3: RLHF — Reinforcement Learning from Human Feedback

SFT teaches the model *what* to say. RLHF teaches it *how good* an answer is — capturing nuance that cannot be expressed with a single 'correct' response.

Why SFT alone is not enough:

Prompt: "Write a Python function to sort a list"

Response A: `def sort_list(lst): return sorted(lst)`

Response B: `def sort_list(lst):
Uses bubble sort - O(n²)
for i in range(len(lst)):
for j in range(len(lst)-i-1):
if lst[j] > lst[j+1]:
lst[j], lst[j+1] = lst[j+1], lst[j]
return lst`

Both are 'correct' — but A is far better (shorter, Pythonic, uses built-in).
A simple SFT loss cannot capture this preference hierarchy.

RLHF in 3 steps:

Step 1 — Human preference data collection

Show human raters two responses to the same prompt. They label: 'Response A is better than B'. Collect millions of (prompt, chosen, rejected) triplets.

Step 2 — Train a Reward Model (RM)

The RM is a separate Transformer that predicts: $RM(prompt, response) \rightarrow \text{scalar score}$. Train it so: $RM(prompt, chosen) > RM(prompt, rejected)$ for all preference pairs.

Step 3 — PPO fine-tuning against the RM

Use Proximal Policy Optimisation (RL algorithm) to fine-tune the LLM. Objective: maximise $RM(prompt, generated_response)$. Constraint: do not drift too far from the SFT model (KL divergence penalty).

Who uses RLHF: Early ChatGPT (InstructGPT paper, OpenAI 2022), GPT-4, Claude (with Constitutional AI variant from Anthropic).

Section 5 — Stage 3 (Modern): DPO — Direct Preference Optimisation

RLHF works but is complex: three models simultaneously, unstable PPO training, massive GPU memory. **DPO (Rafailov et al., 2023) eliminates the reward model entirely.**

The DPO loss function:

```
DPO Loss = -log sigmoid( beta x log[pi(chosen)/pi_ref(chosen)]  
- beta x log[pi(rejected)/pi_ref(rejected)] )
```

Where:

pi = the model being trained

pi_ref = reference model (frozen SFT model)

beta = temperature controlling deviation from reference

sigmoid = sigmoid function

Plain English: DPO directly pushes the model to assign higher probability to the chosen response and lower probability to the rejected response — without ever training a separate reward model. Same result, half the complexity.

DPO vs RLHF comparison:

Property	RLHF	DPO
Reward model needed	Yes — separate training run	No
Training stability	Low — PPO is sensitive to hyperparams	High — standard loss function
GPU memory	3x model size in memory	2x model size in memory
Implementation complexity	Very high	Moderate
Output quality	Slightly higher ceiling	95% of RLHF quality
Who uses it	OpenAI (early), Anthropic	Meta LLaMA-3, Mistral, Zephyr, most OSS

Section 6 — Constitutional AI: Anthropic's Alignment Approach

Anthropic (Claude's creator) uses a variant called **Constitutional AI (CAI)** — worth knowing since Claude powers your AI engineering course:

Step 1 – Generate responses

Model generates responses to potentially harmful or edge-case prompts

Step 2 – Self-critique using a 'Constitution'

The model critiques its own response against ~16 human-written principles:

"Is this response harmful? Is it honest? Does it respect user autonomy?"

Step 3 – Revision

The model rewrites the response to comply with the constitution

Step 4 – Use revised data for RLHF/DPO

The (original, revised) pairs become preference data – no human labelling needed

This is also called **RLAIF (RL from AI Feedback)** and is increasingly used by Google, Meta, and others. It dramatically reduces the amount of expensive human labelling needed while maintaining alignment quality.

Section 7 — The Compute Equation: $C = 6ND$

Every training run can be estimated with one equation:

$$C = 6 \times N \times D$$

C = Training compute (FLOPs)

N = Number of model parameters

D = Number of training tokens

6 = Approximate FLOPs per parameter per token (forward + backward pass)

Real compute numbers for major models:

Model	N (params)	D (tokens)	C (FLOPs)	Est. Cost
GPT-3	175B	300B	3.1×10^{23}	~\$5M
LLaMA-3 8B	8B	15T	7.2×10^{23}	~\$2M
LLaMA-3 70B	70B	15T	6.3×10^{24}	~\$15M
LLaMA-3 405B	405B	15T	3.6×10^{25}	~\$90M
GPT-4 (est.)	~1.8T	~13T	$\sim 1.4 \times 10^{26}$	~\$100M+

Section 8 — Chinchilla Scaling Laws

DeepMind's 2022 Chinchilla paper answered: *for a fixed compute budget, what is the optimal balance between model size and training tokens?*

Chinchilla optimal: $D \approx 20 \times N$

For a 7B parameter model:

Optimal tokens = $20 \times 7B = 140B$ tokens

LLaMA-3 trained 8B model on 15T tokens:

$15T / 8B = 1,875x$ – massively over-trained vs Chinchilla

Reason: cheaper to train once on more data, deploy inference many times

Key finding: GPT-3 (175B params, 300B tokens) was severely under-trained.

A 70B model trained on 1.4T tokens (Chinchilla-optimal) outperformed GPT-3.

Over-training is now standard practice. Train on far more tokens than Chinchilla-optimal because inference is much cheaper than training. A model trained on 15T tokens is smarter per parameter than one trained on 140B tokens, and you only pay the training cost once but run inference millions of times per day.

Section 9 — FAANG Training Strategies at Scale

Meta — LLaMA-3 (Open Source, Published Recipe)

Meta trained LLaMA-3 on 15 trillion tokens across three sizes (8B, 70B, 405B). They used a custom data pipeline to filter CommonCrawl down to high-quality text, deduplication at scale, and 8% code from GitHub (which improves reasoning even on non-code tasks). They published the full training recipe — this is why LLaMA-3 is the foundation for most enterprise fine-tuning today.

Google — Gemini (Native Multimodal Pre-Training)

Google trained Gemini natively multimodal — images, audio, video, and text interleaved during pre-training, not added as an adapter later. This is architecturally different from GPT-4V, which adds vision post-training. Native multimodal pre-training produces better cross-modal reasoning because the model learns joint representations from day one.

OpenAI — GPT-4 (Process Reward Models)

OpenAI's GPT-4 technical report hints at using Process Reward Models (PRMs) during training — rewarding not just the final answer but the intermediate reasoning steps. This is different from outcome-only RLHF. PRMs are the foundation for o1/o3's chain-of-thought reasoning capabilities and represent the frontier of alignment research.

Microsoft — Phi-3 (Data Quality Over Scale)

Microsoft's 3.8B Phi-3-mini was trained on 3.4 trillion tokens of filtered, synthetic 'textbook' data generated by GPT-4. It outperforms models 10x its size on reasoning benchmarks. The insight: synthetic data curated for quality and educational content beats raw web crawl at a fraction of the cost. This makes high-quality models accessible without \$100M budgets.

Section 10 — Architect's Lens: 3 Training Decisions in Production

Decision 1 — Fine-Tune vs Prompt Engineering vs RAG

Before reaching for fine-tuning, ask these questions in order:

Question 1: Can prompt engineering solve this?

→ If yes: use prompting. Zero cost, instant iteration. Start here always.

Question 2: Does the model lack knowledge (not behaviour)?

→ If yes: use RAG (Weeks 3-4). No training needed, knowledge stays fresh.

Question 3: Do you need a consistent style, format, or domain behaviour?

→ If yes: SFT is appropriate. Budget \$10K-\$500K.

Question 4: Do you need safety guardrails or preference alignment?

→ If yes: DPO on top of SFT. Budget additional \$1K-\$50K.

Fine-tuning is overused. Most problems are solved more cheaply with better prompts or RAG.

Decision 2 — Which Base Model to Fine-Tune

Requirement	Recommended Base Model
Open, modifiable, on-premise	LLaMA-3, Mistral, Qwen-2
Best quality, API only	GPT-4o, Claude Sonnet 4
Multilingual (100+ languages)	LLaMA-3, Qwen-2, Aya Expanse
Code-focused	CodeLlama, DeepSeek-Coder, Qwen2.5-Coder
Small / edge / mobile deployment	Phi-3 mini, Gemma-2B, LLaMA-3 1B

Decision 3 — Training Cost Estimation Before Committing

Use $C = 6ND$ to estimate cost before spending:

SFT cost estimate for 10,000 examples x 512 tokens avg:

$D = 10,000 \times 512 = 5.12\text{M tokens}$

$C = 6 \times 7B \times 5.12\text{M} = 2.15 \times 10^{17} \text{ FLOPs}$

On 8 x A100 (312 TFLOPS each): $\sim 2,880 \text{ GPU-seconds} \approx 1 \text{ hour}$

AWS p4d.24xlarge: \$32/hr → $\sim \$32 \text{ per fine-tuning run}$

Scale to 1M examples:

$\sim \$3,200 \text{ per run}$ — still affordable vs \$100M pre-training

Rule: always estimate before you commit. A \$32 SFT run often outperforms a \$100K custom pre-training run when data quality is high.

Section 11 — Hands-On Exercise: SFT Dataset Inspection + Cost Calculator

Inspect a real SFT dataset and calculate fine-tuning cost.

Setup:

```
cd /Users/thaya/AI-Engineering/week-01/notebooks
source /Users/thaya/AI-Engineering/.venv/bin/activate
pip install datasets
```

Create: week-01/notebooks/day4_training_lifecycle.py

```
from datasets import load_dataset

# Load a real SFT dataset
dataset = load_dataset('HuggingFaceH4/ultrachat_200k', split='train_sft[:100]')

print(f'Dataset size (sample): {len(dataset)}')
print(f'Columns: {dataset.column_names}')
print()

# Inspect a few examples
for i, example in enumerate(dataset.select(range(3))):
    messages = example['messages']
    print(f'--- Example {i+1} ---')
    for msg in messages:
        role = msg['role'].upper()
        content = msg['content'][:150]
        print(f' [{role}]: {content}...')
    print()

# Compute training cost estimate
total_tokens = sum(
    len(' '.join(m['content'] for m in ex['messages']).split()) * 1.3
    for ex in dataset
)
print(f'Approx tokens in 100 examples: {int(total_tokens):,}')
total_full = total_tokens * 2000 # scale to full 200K
cost_flops = 6 * 8e9 * total_full
gpu_seconds = cost_flops / 312e12 # A100 TFLOPS
gpu_hours = gpu_seconds / (3600 * 8) # 8xA100
print(f'Estimated GPU hours on 8xA100: ~{gpu_hours:.1f} hrs')
print(f'AWS cost estimate: ~${gpu_hours * 32:.0f}')
```

What You Will Observe:

- Real SFT data is structured as [system, user, assistant] message triplets
- Each example shows exactly the (instruction → response) format the model learns from
- The cost calculation proves SFT is affordable even at 200K examples
- UltraChat data covers diverse topics — this diversity is what makes SFT generalise

Day 4 — Complete Summary

#	Topic	Key Takeaway	Status
1	Three-Stage Pipeline	Pre-Training → SFT → Alignment; most companies skip Stage 1	DONE
2	Pre-Training	Next-token prediction on 15T tokens; learns everything from one simple objective	DONE
3	Training Data	67% web, 13% books, 8% code; quality beats quantity at trillion-token scale	DONE
4	Pre-Training Infrastructure	Meta: 16K H100s, 90 days; GPT-4: 25K A100s, ~\$100M	DONE
5	SFT	Instruction-response pairs; base model becomes assistant; ChatML format	DONE
6	SFT Datasets	OpenAssistant, ShareGPT, Alpaca, FLAN, LLaMA-3 internal 10M examples	DONE
7	RLHF	Human prefs → Reward Model → PPO; 3 models; used by early ChatGPT, Claude	DONE
8	DPO	No reward model; single loss; 2x memory; used by LLaMA-3, Mistral, Zephyr	DONE
9	Constitutional AI	Anthropic RLAIIF; self-critique against 16 principles; fewer human labels	DONE
10	C = 6ND	Training compute equation; GPT-4 ~\$100M; SFT on 10K examples ~\$32	DONE
11	Chinchilla Laws	Optimal D = 20xN; modern models intentionally over-train for inference efficiency	DONE
12	FAANG Training	Meta (LLaMA recipe), Google (multimodal), OpenAI (PRMs), Microsoft (Phi/synthetic)	DONE
13	Architect's Lens	Prompt first, RAG second, fine-tune third; base model selection; C=6ND cost estimation	DONE
14	Hands-On Exercise	SFT dataset inspection + cost calculator — day4_training_lifecycle.py	ACTION

Preview — Day 5: Scaling Laws + Prompt Engineering (Friday, Jul 10)

Now that you understand how models are trained, Day 5 covers two critical practitioner skills: **Scaling Laws** (predicting model capability from compute before training) and **Prompt Engineering** (getting the most out of any

model without changing its weights).

Day 5 will cover:

- Neural scaling laws — loss as a power law of compute, data, and parameters
- Emergent abilities — why capabilities appear suddenly at certain scales
- Prompt engineering techniques: zero-shot, few-shot, chain-of-thought, self-consistency
- System prompts and their role in production AI systems
- Prompt injection risks and basic defences
- Architect's Lens: when to prompt vs fine-tune vs RAG

EAGLE EYE ADDENDUM — Day 4: Critical Training Hyperparameters

Section 12 — Learning Rate, Batch Size, and Checkpointing

The previous sections covered WHAT training stages exist (pre-training, SFT, RLHF/DPO). This addendum covers HOW each stage runs — the three hyperparameters you control that have the largest impact on training quality and stability.

Learning Rate — The Most Important Hyperparameter

The learning rate (lr) controls how large a step the optimizer takes to update weights after each gradient computation. Too large: training diverges (loss goes to NaN). Too small: training converges slowly or gets stuck.

Regime	Typical LR	Effect
Too high	> 1e-3	Loss explodes — NaN weights
Good range	1e-4 to 5e-5	Stable convergence
Too low	< 1e-6	Converges to poor local minimum
SFT range	2e-5 to 1e-4	Standard for fine-tuning 7B-70B models
Pre-training	1e-4 to 3e-4	Meta LLaMA-3 uses 3e-4 with warmup+decay

Learning Rate Schedule: Warmup + Cosine Decay

A constant learning rate is almost never used in practice. The standard LLM training schedule has two phases:

- **Warmup phase** (first 1–5% of steps): lr starts near 0 and linearly increases to the target lr. This prevents early instability when weights are random.
- **Cosine decay phase** (remaining steps): lr decreases from target lr to a minimum (typically 10% of peak) following a cosine curve. This improves final quality.

```
import math

def cosine_lr_schedule(step, total_steps, lr_max, lr_min, warmup_steps):
    if step < warmup_steps:
        return lr_max * (step / warmup_steps)          # linear warmup
    progress = (step - warmup_steps) / (total_steps - warmup_steps)
    return lr_min + 0.5 * (lr_max - lr_min) * (1 + math.cos(math.pi * progress))

# Meta LLaMA-3 8B training schedule:
# lr_max=3e-4, lr_min=3e-5, warmup=2000 steps
# Total training ~37,500 steps (15T tokens / 4096 ctx / batch 100)
```


FAANG EXAMPLE — Meta LLaMA-3: warmup for 2,000 steps (0.005% of training), then cosine decay to $lr_{min} = lr_{max} / 10$. Google's Gemini also uses warmup+cosine. This schedule is so standard it is the default in every major training framework (PyTorch Lightning, Hugging Face Trainer, DeepSpeed).

Batch Size — Throughput vs Gradient Quality

Batch size is the number of training examples processed before a single weight update. For LLM training, 'batch size' means something more nuanced than a simple number:

- **Per-GPU micro-batch size:** how many sequences fit in one GPU's memory at once. Limited by VRAM. For a 7B model on A100 80GB, typically 1–4.
- **Gradient accumulation steps:** run N micro-batches, accumulate gradients, then do one update. Simulates a larger batch size without extra VRAM.
- **Global batch size** = per-GPU micro-batch × gradient accumulation × num GPUs. This is what the learning rate is calibrated to.

```
# Real numbers from Meta LLaMA-3 8B training:
# 16,384 H100 GPUs
# Per-GPU micro-batch: 1 sequence of 8,192 tokens
# No gradient accumulation (each GPU gets full batch)
# Global batch size: 16,384 sequences × 8,192 tokens
#                               = 134M tokens per step

# For SFT on a single 8xA100 machine:
per_gpu_batch = 2           # sequences per GPU
gradient_accum = 4          # accumulate 4 micro-batches
num_gpus = 8
global_batch = per_gpu_batch * gradient_accum * num_gpus # = 64

# Linear learning rate scaling rule:
# If global_batch doubles, scale lr proportionally
# lr = base_lr × (global_batch / reference_batch)
```

RULE OF THUMB: larger global batch size = smoother gradient estimates = can use higher learning rate. Halve the batch size → halve the learning rate (linear scaling rule, Goyal et al. 2017, Facebook AI). LLaMA-3 uses a global batch of 134M tokens/step; typical SFT uses 1M tokens/step.

Checkpointing — Saving Model State During Long Runs

Pre-training runs for 30–90 days on thousands of GPUs. Hardware failures, power outages, and OOM errors are inevitable. Checkpointing saves the complete model state at regular intervals so training can resume without starting over.

What a checkpoint contains:

- Model weights (the transformer parameters at this step)
- Optimizer state (Adam's running mean m and variance v for every parameter) — this is 2× the size of the model weights
- Training step number, learning rate schedule position
- Random number generator state (for reproducibility of subsequent steps)

```
# Checkpoint size reality:
# LLaMA-3 70B model weights (FP16): 140 GB
# + Adam optimizer state (FP32): +560 GB (2 tensors per param, FP32)
# + activations/misc: +~50 GB
# Total checkpoint: ~750 GB
#
# Meta saves a checkpoint every 1,000 steps (~6 hours)
# They keep last 5 checkpoints, delete older ones
# Total checkpoint storage: ~3.75 TB

# In HuggingFace Trainer:
from transformers import TrainingArguments

args = TrainingArguments(
    output_dir="./checkpoints",
    save_steps=500,          # checkpoint every 500 steps
    save_total_limit=3,      # keep only last 3 checkpoints
    load_best_model_at_end=True, # load checkpoint with best eval loss
    evaluation_strategy="steps",
    eval_steps=500,
)
```

ARCHITECT NOTE: Never start a multi-day training run without checkpointing. A 90-day pre-training run without checkpoints loses \$1M+ if a datacenter failure occurs on day 89. Meta checkpoints LLaMA every ~6 hours. For SFT runs, checkpoint every 500–1,000 steps. Checkpoint storage cost is trivial compared to retraining cost.

Hyperparameter Quick Reference

Hyperparameter	Typical Values	What Happens if Wrong
Learning rate (peak)	3e-4 (pre-train), 2e-5 (SFT)	Too high: NaN loss. Too low: no convergence.
LR warmup steps	1–5% of total steps	Skip warmup: unstable early training, possible divergence.
LR decay	Cosine to $lr_min = lr_max/10$	Constant LR: 1–3% worse final quality vs cosine.
Batch size (global)	1M–100M tokens/step	Too small: noisy gradients. Too large: diminishing returns.
Checkpoint frequency	Every 500–1,000 steps	Too rare: lose hours/days of compute on failure.

EAGLE EYE ADDENDUM — DAY 4 | Epoch, Overfitting, Validation — Training Fundamentals

Day 4 covered the three-stage training lifecycle (pre-training → SFT → alignment) and learning rate schedules. Four foundational training concepts were missing: epoch, validation set, overfitting, and early stopping. Every AI engineer working on fine-tuning must understand these.

Epoch, Step, and Iteration

```
# Definitions:
# epoch      = one full pass through the entire training dataset
# step       = one gradient update on one batch
# iteration  = same as step; used interchangeably

# Relationship:
# steps_per_epoch = dataset_size / batch_size
# total_steps = epochs * steps_per_epoch

# Example: 100K training examples, batch=32, 3 epochs
dataset_size = 100_000
batch_size   = 32
epochs       = 3
steps_per_epoch = dataset_size // batch_size      # 3,125 steps
total_steps    = epochs * steps_per_epoch         # 9,375 steps

# In Hugging Face:
from transformers import TrainingArguments
args = TrainingArguments(
    num_train_epochs=3,          # epochs – most common to set
    per_device_train_batch_size=32,
    max_steps=-1,               # -1 = use num_train_epochs; positive overrides
    # OR set max_steps=9375 instead of num_train_epochs
)

# Fine-tuning rule of thumb:
# LoRA fine-tuning: 1-5 epochs (dataset typically small, 1K-50K examples)
# Full fine-tuning: 1-3 epochs (larger dataset; more compute-expensive)
# SFT on pre-trained: 1-2 epochs (risk of overfitting rises fast after epoch 2)
```

Train / Validation / Test Split

Every training run needs three distinct data splits. Mixing these is one of the most common mistakes in AI engineering — and it leads to models that look great in training but fail in production.

```
# Standard splits for LLM fine-tuning:
# Train:      80-90% — model sees this during training
# Validation: 5-10% — model NEVER sees; used to tune hyperparameters and detect overfitting
# Test:       5-10% — held out until FINAL evaluation; touched once, at the end

from datasets import load_dataset
```

```

dataset = load_dataset("your/sql-dataset")
split = dataset["train"].train_test_split(test_size=0.15, seed=42)
train_data = split["train"]    # 85%
temp      = split["test"].train_test_split(test_size=0.5, seed=42)
val_data  = temp["train"]      # 7.5%
test_data = temp["test"]       # 7.5%

# CRITICAL RULES:
# Never use test data to pick hyperparameters – you'll overfit to the test set
# Never shuffle after splitting – you'll leak data across splits
# For SQL fine-tuning: ensure same database does NOT appear in both train and val
# (schema-level leakage is subtle and kills generalisation)

```

Overfitting — What It Is and How to Detect It

Overfitting is when a model memorises the training data instead of learning the underlying pattern. It performs well on training examples it has seen, but fails on new, unseen examples. It is the most common failure mode in fine-tuning.

```

# Detecting overfitting: training loss vs validation loss divergence
#
# Epoch  Train Loss  Val Loss  Status
#   1      2.40      2.45    Normal – both decreasing
#   2      1.80      1.85    Normal
#   3      1.30      1.40    Good generalisation
#   4      0.90      1.55    OVERFITTING – val loss rising while train loss falls
#   5      0.60      1.90    Severe overfitting
#
# The validation loss is the signal. Train loss ALWAYS decreases with more training.

# Configure validation logging with Hugging Face:
from transformers import TrainingArguments, Trainer

args = TrainingArguments(
    evaluation_strategy="steps",          # evaluate every N steps
    eval_steps=500,                      # run validation every 500 steps
    logging_steps=100,                   # log train loss every 100 steps
    load_best_model_at_end=True,         # keep best checkpoint, not latest
    metric_for_best_model="eval_loss",   # what "best" means
    greater_is_better=False,             # lower eval_loss = better
    save_strategy="steps",
    save_steps=500,
)

trainer = Trainer(
    model=model,
    args=args,
    train_dataset=train_data,
    eval_dataset=val_data,               # validation set – NOT test set
)
trainer.train()

```

Early Stopping

Early stopping automatically halts training when validation loss stops improving — saving compute and preventing overfitting. It is the standard practice for all LLM fine-tuning.

```
from transformers import EarlyStoppingCallback

trainer = Trainer(
    model=model,
    args=args,
    train_dataset=train_data,
    eval_dataset=val_data,
    callbacks=[
        EarlyStoppingCallback(
            early_stopping_patience=3,    # stop if val loss doesn't improve for 3 evals
            early_stopping_threshold=0.0, # minimum improvement to count as improvement
        )
    ],
)
trainer.train()

# patience=3 means: if val_loss at eval step N, N+1, N+2 are all >= val_loss at N-1,
# stop training. The checkpoint with the BEST val_loss is saved (load_best_model_at_end=True).

# Practical settings for LLM fine-tuning:
# LoRA (1K-10K examples):  patience=3, eval every 50-100 steps
# Full FT (50K+ examples): patience=3, eval every 200-500 steps
# SFT on large dataset:    patience=2, eval every 500 steps (compute expensive)
```

Overfitting Prevention Strategies

Technique	How It Works	When to Apply
Weight decay (L2)	Penalises large weights; keeps model generalised	Always; typical 0.01-0.1
Dropout	Randomly zeros activations during training; forces redundancy	SFT; not for pre-training
Data augmentation	Paraphrase questions; swap table aliases in SQL	When dataset < 5K examples
Early stopping	Stop when validation loss stops improving	All fine-tuning runs
Smaller LR	Lower learning rate = slower memorisation	When dataset is small
Data mixing	Mix domain data + general data (90/10 or 95/5)	When forgetting is a risk (Days 9-10)

FAANG RULE: At Meta, every fine-tuning run logs train loss and eval loss to an internal dashboard (Manifold). If eval loss diverges from train loss by more than 15% for 3 consecutive evaluations, an alert fires and the run is paused. This is early stopping as an infrastructure concern, not just a training decision. At Amazon, SageMaker Experiments automatically tracks best checkpoint by val loss across all hyperparameter tuning jobs — picking the best model is always based on val loss, never train loss.

EAGLE EYE ADDENDUM — Day 4: Backpropagation — How Gradients Flow Through the Network

The Forward and Backward Pass

Training a neural network is a two-step loop. The forward pass computes the model's prediction and measures error (loss). The backward pass propagates that error signal backwards through every layer using the chain rule of calculus, computing the gradient of the loss with respect to every weight. The optimizer then adjusts each weight by a small step in the opposite direction of its gradient.

```
# Minimal transformer training loop – forward + backward pass
import torch

model      = MyTransformer(vocab_size=32000, d_model=512)
optimizer  = torch.optim.AdamW(model.parameters(), lr=3e-4)
loss_fn    = torch.nn.CrossEntropyLoss()

for batch in dataloader:
    input_ids = batch["input_ids"]          # (B, T)
    labels    = batch["labels"]             # (B, T) – next-token targets

    # ■■ FORWARD PASS ████████████████████████████████████
    # 1. Each layer transforms the hidden states
    # 2. Final linear head projects to vocab logits
    # 3. Loss measures how wrong the predictions are
    logits = model(input_ids)               # (B, T, V)
    loss   = loss_fn(logits.view(-1, logits.size(-1)), labels.view(-1))

    # ■■ BACKWARD PASS (backpropagation) ████████████████
    # 1. Compute gradient of loss w.r.t. every parameter
    #     using the chain rule:  $dL/dw = dL/dout * dout/dw$ 
    # 2. PyTorch traces the computation graph built during forward pass
    # 3. .backward() fills .grad for every parameter
    optimizer.zero_grad()                  # clear previous gradients
    loss.backward()                         # ← backpropagation happens here

    # ■■ OPTIMIZER STEP ████████████████████████████████████
    # AdamW:  $W = W - lr * (\hat{m} / (\sqrt{\hat{v}} + \epsilon)) - lr * wd * W$ 
    torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
    optimizer.step()

print(f"Loss: {loss.item():.4f}")
```

The Chain Rule — Why Backprop Works

```
# Chain rule: dL/dw1 = dL/dw2 * dw2/dw1 (for any two layers 1, 2)
#
# For a 3-layer network: L = loss(h3)  h3 = w3 * h2  h2 = w2 * h1  h1 = w1 * x
#
# dL/dw1 = dL/dh3 * dh3/dh2 * dh2/dh1 * dh1/dw1
#           (1 scalar) (w3)         (w2)         (h1)
```



```
#
# PyTorch autograd builds a computation graph during forward pass.
# .backward() traverses the graph in reverse, multiplying Jacobians.

# PRACTICAL IMPLICATION:
# Every forward-pass operation must be differentiable (or have a surrogate gradient).
# Activation functions: ReLU, GELU, SiLU — all differentiable.
# Discrete operations (argmax, sampling) — NOT differentiable → can't backprop through them.
# This is why LLM training uses teacher-forcing (feed ground-truth tokens as inputs)
# rather than the model's own predicted tokens.

# MEMORY: backward pass stores ALL intermediate activations to compute gradients.
# An 8B LLM forward pass: ~16 GB (weights)
# Same model training (full backprop): ~80-100 GB (weights + activations + gradients + optimizer state)
# ← This is why training requires 4-8x more GPU memory than inference
```

Pass	Direction	What Happens	Memory Use
Forward	Input → Output	Compute logits + loss; build computation graph	Activations stored for backward
Backward	Output → Input	Chain rule: compute dL/dW for every parameter	Uses stored activations; allocates gradient tensors
Optimizer	—	$W = W - lr * dL/dW$ (Adam adds momentum terms)	Optimizer state: 2x params (momentum + variance)

WHY IT MATTERS FOR FINE-TUNING: Full fine-tuning backpropagates through ALL layers. LoRA only backpropagates through the low-rank adapter matrices A and B — all other layers are frozen so their gradients are not computed or stored. This is why LoRA is 5-10× more memory-efficient during training: gradient memory is proportional to the number of trainable parameters.